

yxq76abae

January 31, 2024

Regression is a type of supervised machine learning technique used for predicting a continuous target variable based on one or more input features. The goal of regression is to find a mathematical relationship (model) that maps the input features to the target variable. There are several regression algorithms you can use, depending on the nature of your data and the problem you are trying to solve. Here's an overview of some common regression techniques:

1. **Linear Regression:**

- Simple linear regression models the relationship between a single input feature and the target variable as a straight line.
- Multiple linear regression extends this to multiple input features.
- The model assumes a linear relationship between the features and the target.

2. **Support Vector Regression (SVR):**

- SVR is a regression technique based on support vector machines.
- It seeks to find a hyperplane that best fits the data, allowing a margin of error.

3. **K-Nearest Neighbors (KNN)** can be used for regression tasks as well as classification. When applied to regression, it's often referred to as "K-Nearest Neighbors Regression" or "KNN Regression." Instead of predicting a class label, KNN regression predicts a continuous target variable

4. **Polynomial Regression:**

- Polynomial regression extends linear regression by fitting a polynomial equation to the data, allowing it to capture more complex relationships.
- It can model curved relationships, which a simple linear regression cannot.

5. **Decision Tree Regression:**

- Decision tree regression uses decision trees to partition the data into regions and predict the target variable based on the average of data points in each region.

```
[2]: #import library
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

```
[3]: # Data To Dataframe
df=pd.read_csv("CarPrice_Assignment.csv")
```

1 Exploratory data analysis Python

```
[4]: # Shallow copy
df_copy=df.copy
```

```
[5]: # The shape of a DataFrame.
df.shape
```

```
[5]: (205, 26)
```

```
[6]: # column labels of the Dataframe
df.columns
```

```
[6]: Index(['car_ID', 'symboling', 'CarName', 'fueltype', 'aspiration',
        'doornumber', 'carbody', 'drivewheel', 'enginelocation', 'wheelbase',
        'carlength', 'carwidth', 'carheight', 'curbweight', 'enginetype',
        'cylindernumber', 'enginesize', 'fuelsystem', 'boreratio', 'stroke',
        'compressionratio', 'horsepower', 'peakrpm', 'citympg', 'highwaympg',
        'price'],
        dtype='object')
```

```
[7]: # Returns a specified number of rows, string from the top.
df.head()
```

```
[7]:
```

	car_ID	symboling	CarName	fueltype	aspiration	doornumber	\
0	1	3	alfa-romero giulia	gas	std	two	
1	2	3	alfa-romero stelvio	gas	std	two	
2	3	1	alfa-romero Quadrifoglio	gas	std	two	
3	4	2	audi 100 ls	gas	std	four	
4	5	2	audi 100ls	gas	std	four	

	carbody	drivewheel	enginelocation	wheelbase	...	enginesize	\
0	convertible	rwd	front	88.6	...	130	
1	convertible	rwd	front	88.6	...	130	
2	hatchback	rwd	front	94.5	...	152	
3	sedan	fwd	front	99.8	...	109	
4	sedan	4wd	front	99.4	...	136	

	fuelsystem	boreratio	stroke	compressionratio	horsepower	peakrpm	citympg	\
0	mpfi	3.47	2.68	9.0	111	5000	21	
1	mpfi	3.47	2.68	9.0	111	5000	21	
2	mpfi	2.68	3.47	9.0	154	5000	19	
3	mpfi	3.19	3.40	10.0	102	5500	24	
4	mpfi	3.19	3.40	8.0	115	5500	18	

	highwaympg	price
0	27	13495.0

```

1      27  16500.0
2      26  16500.0
3      30  13950.0
4      22  17450.0

```

[5 rows x 26 columns]

```

[8]: # Return the last 5 rows
df.tail()

```

```

[8]:   car_ID  symboling      CarName fueltype aspiration doornumber \
200    201        -1  volvo 145e (sw)      gas          std         four
201    202        -1    volvo 144ea      gas         turbo         four
202    203        -1    volvo 244dl      gas          std         four
203    204        -1      volvo 246  diesel         turbo         four
204    205        -1    volvo 264gl      gas         turbo         four

      carbody drivewheel enginelocation  wheelbase  ...  enginesize  fuelsystem \
200    sedan         rwd         front     109.1  ...        141        mpfi
201    sedan         rwd         front     109.1  ...        141        mpfi
202    sedan         rwd         front     109.1  ...        173        mpfi
203    sedan         rwd         front     109.1  ...        145         idi
204    sedan         rwd         front     109.1  ...        141        mpfi

      boreratio  stroke  compressionratio  horsepower  peakrpm  citympg  \
200         3.78    3.15              9.5          114     5400        23
201         3.78    3.15              8.7          160     5300        19
202         3.58    2.87              8.8          134     5500        18
203         3.01    3.40             23.0          106     4800        26
204         3.78    3.15              9.5          114     5400        19

      highwaympg  price
200           28  16845.0
201           25  19045.0
202           23  21485.0
203           27  22470.0
204           25  22625.0

```

[5 rows x 26 columns]

```

[9]: # specific rows of a DataFrame ( "integer location" Method)
df.iloc[100:200]

```

```

[9]:   car_ID  symboling      CarName fueltype aspiration doornumber \
100    101          0  nissan nv200      gas          std         four
101    102          0   nissan dayz      gas          std         four
102    103          0   nissan fuga      gas          std         four

```

103	104	0	nissan otti	gas	std	four
104	105	3	nissan teana	gas	std	two
..
195	196	-1	volvo 144ea	gas	std	four
196	197	-2	volvo 244dl	gas	std	four
197	198	-1	volvo 245	gas	std	four
198	199	-2	volvo 264gl	gas	turbo	four
199	200	-1	volvo diesel	gas	turbo	four

	carbody	drivewheel	enginelocation	wheelbase	...	enginesize	\
100	sedan	fwd	front	97.2	...	120	
101	sedan	fwd	front	100.4	...	181	
102	wagon	fwd	front	100.4	...	181	
103	sedan	fwd	front	100.4	...	181	
104	hatchback	rwd	front	91.3	...	181	
..	
195	wagon	rwd	front	104.3	...	141	
196	sedan	rwd	front	104.3	...	141	
197	wagon	rwd	front	104.3	...	141	
198	sedan	rwd	front	104.3	...	130	
199	wagon	rwd	front	104.3	...	130	

	fuelsystem	boreratio	stroke	compressionratio	horsepower	peakrpm	\
100	2bbl	3.33	3.47	8.5	97	5200	
101	mpfi	3.43	3.27	9.0	152	5200	
102	mpfi	3.43	3.27	9.0	152	5200	
103	mpfi	3.43	3.27	9.0	152	5200	
104	mpfi	3.43	3.27	9.0	160	5200	
..	
195	mpfi	3.78	3.15	9.5	114	5400	
196	mpfi	3.78	3.15	9.5	114	5400	
197	mpfi	3.78	3.15	9.5	114	5400	
198	mpfi	3.62	3.15	7.5	162	5100	
199	mpfi	3.62	3.15	7.5	162	5100	

	citympg	highwaympg	price
100	27	34	9549.0
101	17	22	13499.0
102	17	22	14399.0
103	19	25	13499.0
104	19	25	17199.0
..
195	23	28	13415.0
196	24	28	15985.0
197	24	28	16515.0
198	17	22	18420.0
199	17	22	18950.0

[100 rows x 26 columns]

```
[10]: # describe() method gives us summary statistics for numerical columns in df
      ↪ DataFrame.
      df.describe().T
```

```
[10]:
```

	count	mean	std	min	25%	\
car_ID	205.0	103.000000	59.322565	1.00	52.00	
symboling	205.0	0.834146	1.245307	-2.00	0.00	
wheelbase	205.0	98.756585	6.021776	86.60	94.50	
carlength	205.0	174.049268	12.337289	141.10	166.30	
carwidth	205.0	65.907805	2.145204	60.30	64.10	
carheight	205.0	53.724878	2.443522	47.80	52.00	
curbweight	205.0	2555.565854	520.680204	1488.00	2145.00	
enginesize	205.0	126.907317	41.642693	61.00	97.00	
boreratio	205.0	3.329756	0.270844	2.54	3.15	
stroke	205.0	3.255415	0.313597	2.07	3.11	
compressionratio	205.0	10.142537	3.972040	7.00	8.60	
horsepower	205.0	104.117073	39.544167	48.00	70.00	
peakrpm	205.0	5125.121951	476.985643	4150.00	4800.00	
citympg	205.0	25.219512	6.542142	13.00	19.00	
highwaympg	205.0	30.751220	6.886443	16.00	25.00	
price	205.0	13276.710571	7988.852332	5118.00	7788.00	

	50%	75%	max
car_ID	103.00	154.00	205.00
symboling	1.00	2.00	3.00
wheelbase	97.00	102.40	120.90
carlength	173.20	183.10	208.10
carwidth	65.50	66.90	72.30
carheight	54.10	55.50	59.80
curbweight	2414.00	2935.00	4066.00
enginesize	120.00	141.00	326.00
boreratio	3.31	3.58	3.94
stroke	3.29	3.41	4.17
compressionratio	9.00	9.40	23.00
horsepower	95.00	116.00	288.00
peakrpm	5200.00	5500.00	6600.00
citympg	24.00	30.00	49.00
highwaympg	30.00	34.00	54.00
price	10295.00	16503.00	45400.00

```
[11]: # prints information about the DataFrame. [number of columns, column labels,
      ↪ column data types, memory usage, range index,]
      df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
#   Column                Non-Null Count  Dtype
---  -
0   car_ID                205 non-null   int64
1   symboling              205 non-null   int64
2   CarName               205 non-null   object
3   fueltype              205 non-null   object
4   aspiration            205 non-null   object
5   doornumber            205 non-null   object
6   carbody               205 non-null   object
7   drivewheel           205 non-null   object
8   enginelocation        205 non-null   object
9   wheelbase             205 non-null   float64
10  carlength             205 non-null   float64
11  carwidth              205 non-null   float64
12  carheight             205 non-null   float64
13  curbweight            205 non-null   int64
14  enginetype            205 non-null   object
15  cylindernumber        205 non-null   object
16  enginesize            205 non-null   int64
17  fuelsystem            205 non-null   object
18  boreratio             205 non-null   float64
19  stroke                205 non-null   float64
20  compressionratio      205 non-null   float64
21  horsepower            205 non-null   int64
22  peakrpm              205 non-null   int64
23  citympg               205 non-null   int64
24  highwaympg           205 non-null   int64
25  price                 205 non-null   float64
dtypes: float64(8), int64(8), object(10)
memory usage: 41.8+ KB

```

```

[12]: # Including only string columns in a DataFrame description
df.describe(include=object)

```

```

[12]:
      CarName fueltype aspiration doornumber carbody drivewheel \
count      205      205      205      205      205      205
unique     147        2        2        2        5        3
top   toyota corona    gas      std      four    sedan    fwd
freq         6     185     168     115     96     120

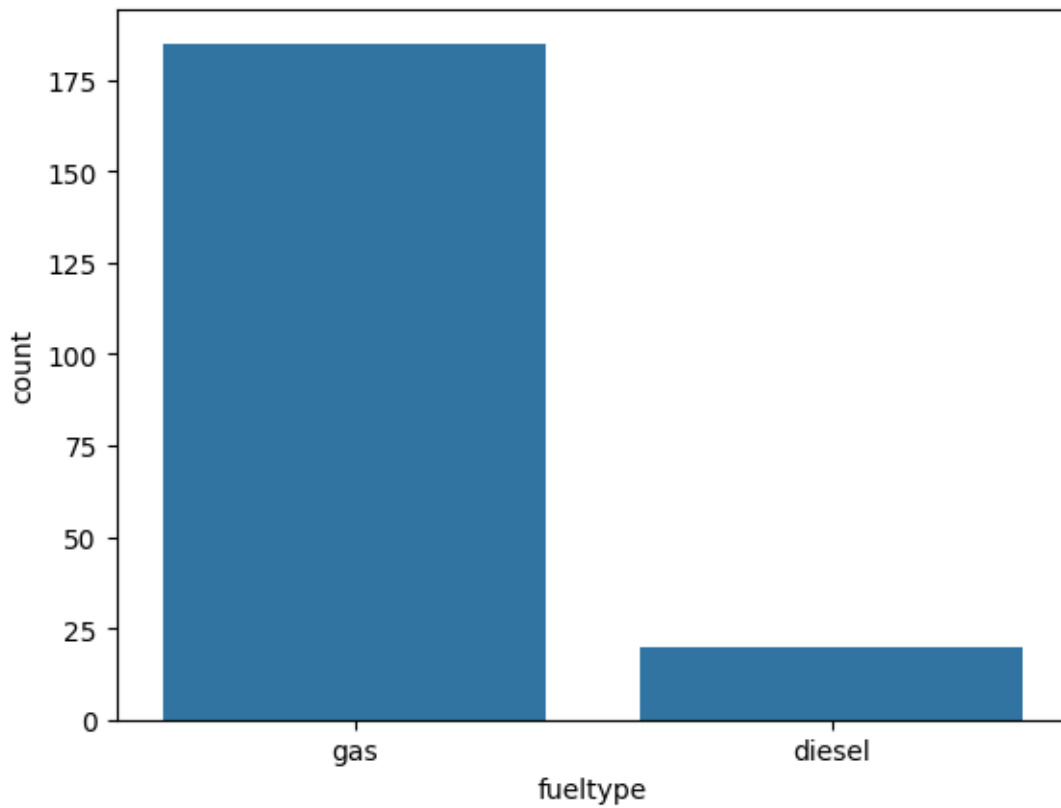
      enginelocation enginetype cylindernumber fuelsystem
count      205      205      205      205
unique        2        7        7        8
top         front      ohc      four    mpfi

```

freq	202	148	159	94
------	-----	-----	-----	----

```
[13]: sns.countplot(data=df,x='fueltype')
```

```
[13]: <Axes: xlabel='fueltype', ylabel='count'>
```



#DATA CLEANING

```
[14]: # To check for duplicate values in a DataFrame
df.duplicated().sum()
```

```
[14]: 0
```

```
[15]: # Missing values in the dataset.
df.isnull().sum()
```

```
[15]: car_ID      0
      symboling  0
      CarName   0
      fueltype  0
      aspiration 0
```

doornumber	0
carbody	0
drivewheel	0
enginelocation	0
wheelbase	0
carlength	0
carwidth	0
carheight	0
curbweight	0
enginetype	0
cylindernumber	0
enginesize	0
fuelsystem	0
boreratio	0
stroke	0
compressionratio	0
horsepower	0
peakrpm	0
citympg	0
highwaympg	0
price	0
dtype:	int64

```
[16]: # index row columns remove
df=df.drop(["car_ID"],axis=1)
```

```
[17]: df.dtypes
```

```
[17]: symboling          int64
CarName              object
fueltype            object
aspiration          object
doornumber          object
carbody             object
drivewheel          object
enginelocation      object
wheelbase          float64
carlength          float64
carwidth           float64
carheight          float64
curbweight          int64
enginetype          object
cylindernumber      object
enginesize          int64
fuelsystem          object
boreratio           float64
stroke             float64
```



```

compressionratio    float64
horsepower           int64
peakrpm              int64
citympg              int64
highwaympg           int64
price                float64
dtype: object

```

```
[18]: df['CarName'].unique()
```

```

[18]: array(['alfa-romero giulia', 'alfa-romero stelvio',
            'alfa-romero Quadrifoglio', 'audi 100 ls', 'audi 100ls',
            'audi fox', 'audi 5000', 'audi 4000', 'audi 5000s (diesel)',
            'bmw 320i', 'bmw x1', 'bmw x3', 'bmw z4', 'bmw x4', 'bmw x5',
            'chevrolet impala', 'chevrolet monte carlo', 'chevrolet vega 2300',
            'dodge rampage', 'dodge challenger se', 'dodge d200',
            'dodge monaco (sw)', 'dodge colt hardtop', 'dodge colt (sw)',
            'dodge coronet custom', 'dodge dart custom',
            'dodge coronet custom (sw)', 'honda civic', 'honda civic cvcc',
            'honda accord cvcc', 'honda accord lx', 'honda civic 1500 gl',
            'honda accord', 'honda civic 1300', 'honda prelude',
            'honda civic (auto)', 'isuzu MU-X', 'isuzu D-Max ',
            'isuzu D-Max V-Cross', 'jaguar xj', 'jaguar xf', 'jaguar xk',
            'maxda rx3', 'maxda glc deluxe', 'mazda rx2 coupe', 'mazda rx-4',
            'mazda glc deluxe', 'mazda 626', 'mazda glc', 'mazda rx-7 gs',
            'mazda glc 4', 'mazda glc custom l', 'mazda glc custom',
            'buick electra 225 custom', 'buick century luxus (sw)',
            'buick century', 'buick skyhawk', 'buick opel isuzu deluxe',
            'buick skylark', 'buick century special',
            'buick regal sport coupe (turbo)', 'mercury cougar',
            'mitsubishi mirage', 'mitsubishi lancer', 'mitsubishi outlander',
            'mitsubishi g4', 'mitsubishi mirage g4', 'mitsubishi montero',
            'mitsubishi pajero', 'Nissan versa', 'nissan gt-r', 'nissan rogue',
            'nissan latio', 'nissan titan', 'nissan leaf', 'nissan juke',
            'nissan note', 'nissan clipper', 'nissan nv200', 'nissan dayz',
            'nissan fuga', 'nissan otti', 'nissan teana', 'nissan kicks',
            'peugeot 504', 'peugeot 304', 'peugeot 504 (sw)', 'peugeot 604sl',
            'peugeot 505s turbo diesel', 'plymouth fury iii',
            'plymouth cricket', 'plymouth satellite custom (sw)',
            'plymouth fury gran sedan', 'plymouth valiant', 'plymouth duster',
            'porsche macan', 'porcshce panamera', 'porsche cayenne',
            'porsche boxter', 'renault 12tl', 'renault 5 gtl', 'saab 99e',
            'saab 99le', 'saab 99gle', 'subaru', 'subaru dl', 'subaru brz',
            'subaru baja', 'subaru r1', 'subaru r2', 'subaru trezia',
            'subaru tribeca', 'toyota corona mark ii', 'toyota corona',
            'toyota corolla 1200', 'toyota corona hardtop',
            'toyota corolla 1600 (sw)', 'toyota carina', 'toyota mark ii',

```

```
'toyota corolla', 'toyota corolla liftback',
'toyota celica gt liftback', 'toyota corolla tercel',
'toyota corona liftback', 'toyota starlet', 'toyota tercel',
'toyota cressida', 'toyota celica gt', 'toyouta tercel',
'volkswagen rabbit', 'volkswagen 1131 deluxe sedan',
'volkswagen model 111', 'volkswagen type 3', 'volkswagen 411 (sw)',
'volkswagen super beetle', 'volkswagen dasher', 'vw dasher',
'vw rabbit', 'volkswagen rabbit', 'volkswagen rabbit custom',
'volvo 145e (sw)', 'volvo 144ea', 'volvo 244dl', 'volvo 245',
'volvo 264gl', 'volvo diesel', 'volvo 246'], dtype=object)
```

```
[19]: df.insert(1, 'Company', df['CarName'].str.split(' ').str[0])
      #df.insert(1, 'company', df['CarName'].str.split(' ').str[0]) #error
```

```
[20]: df.drop(columns=['CarName'], inplace=True)
```

```
[21]: df
```

```
[21]:
```

	symboling	Company	fueltype	aspiration	doornumber	carbody	\
0	3	alfa-romero	gas	std	two	convertible	
1	3	alfa-romero	gas	std	two	convertible	
2	1	alfa-romero	gas	std	two	hatchback	
3	2	audi	gas	std	four	sedan	
4	2	audi	gas	std	four	sedan	
..	
200	-1	volvo	gas	std	four	sedan	
201	-1	volvo	gas	turbo	four	sedan	
202	-1	volvo	gas	std	four	sedan	
203	-1	volvo	diesel	turbo	four	sedan	
204	-1	volvo	gas	turbo	four	sedan	

	drivewheel	engine location	wheelbase	carlength	...	engine size	\
0	rwd	front	88.6	168.8	...	130	
1	rwd	front	88.6	168.8	...	130	
2	rwd	front	94.5	171.2	...	152	
3	fwd	front	99.8	176.6	...	109	
4	4wd	front	99.4	176.6	...	136	
..	
200	rwd	front	109.1	188.8	...	141	
201	rwd	front	109.1	188.8	...	141	
202	rwd	front	109.1	188.8	...	173	
203	rwd	front	109.1	188.8	...	145	
204	rwd	front	109.1	188.8	...	141	

	fuelsystem	bore ratio	stroke	compression ratio	horsepower	peakrpm	\
0	mpfi	3.47	2.68	9.0	111	5000	
1	mpfi	3.47	2.68	9.0	111	5000	

2	mpfi	2.68	3.47	9.0	154	5000
3	mpfi	3.19	3.40	10.0	102	5500
4	mpfi	3.19	3.40	8.0	115	5500
..
200	mpfi	3.78	3.15	9.5	114	5400
201	mpfi	3.78	3.15	8.7	160	5300
202	mpfi	3.58	2.87	8.8	134	5500
203	idi	3.01	3.40	23.0	106	4800
204	mpfi	3.78	3.15	9.5	114	5400

	citympg	highwaympg	price
0	21	27	13495.0
1	21	27	16500.0
2	19	26	16500.0
3	24	30	13950.0
4	18	22	17450.0
..
200	23	28	16845.0
201	19	25	19045.0
202	18	23	21485.0
203	26	27	22470.0
204	19	25	22625.0

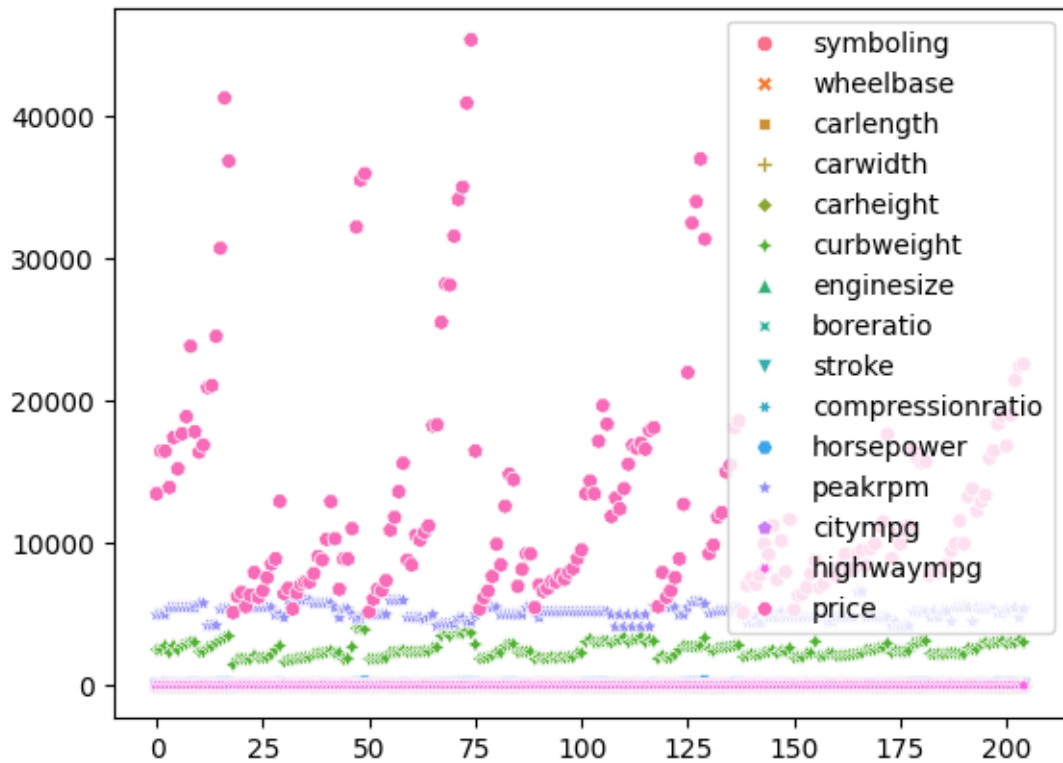
[205 rows x 25 columns]

```
[22]: df['Company'].unique()
```

```
[22]: array(['alfa-romero', 'audi', 'bmw', 'chevrolet', 'dodge', 'honda',
        'isuzu', 'jaguar', 'maxda', 'mazda', 'buick', 'mercury',
        'mitsubishi', 'Nissan', 'nissan', 'peugeot', 'plymouth', 'porsche',
        'porcshce', 'renault', 'saab', 'subaru', 'toyota', 'toyouta',
        'vokswagen', 'volkswagen', 'vw', 'volvo'], dtype=object)
```


```
[23]: sns.scatterplot(df)
```

```
[23]: <Axes: >
```



```
[24]: plt.subplots(figsize=(25,10))
      sns.barplot(data=df,x='Company',y='price')
```

```
[24]: <Axes: xlabel='Company', ylabel='price'>
```



```
[25]: # To select strings you must use the object dtype, but note that this will
      ↪return all object dtype columns
      df.select_dtypes(include=object).columns.to_list()
      #df.select_dtypes(include=object).columns.tolist()
```

```
[25]: ['Company',
      'fueltype',
      'aspiration',
      'doornumber',
      'carbody',
      'drivewheel',
      'enginelocation',
      'enginetype',
      'cylindernumber',
      'fuelsystem']
```

```
[26]: # unues columns
      df.drop(columns=['symboling', 'wheelbase'], inplace=True)
```

```
[27]: # Numerical columns list
      numerical_columns = df.select_dtypes(include=['int64', 'float64']).columns.
      ↪tolist()
      numerical_columns.remove('price')
```

```
[28]: # Label encoding
```

```
[29]: from sklearn.preprocessing import LabelEncoder
      lb=LabelEncoder()
      df['Company']=lb.fit_transform(df['Company'])
      df['fueltype']=lb.fit_transform(df['fueltype'])
      df['aspiration']=lb.fit_transform(df['aspiration'])
      df['doornumber']=lb.fit_transform(df['doornumber'])
      df['carbody']=lb.fit_transform(df['carbody'])
      df['drivewheel']=lb.fit_transform(df['drivewheel'])
      df['enginelocation']=lb.fit_transform(df['enginelocation'])
      df['enginetype']=lb.fit_transform(df['enginetype'])
      df['cylindernumber']=lb.fit_transform(df['cylindernumber'])
      df['fuelsystem']=lb.fit_transform(df['fuelsystem'])
```

```
[30]: numerical_columns
```

```
[30]: ['carlength',
      'carwidth',
      'carheight',
      'curbweight',
      'enginesize',
      'boreratio',
```

```
'stroke',
'compressionratio',
'horsepower',
'peakrpm',
'citympg',
'highwaympg']
```

```
[31]: from sklearn.preprocessing import StandardScaler
ssr=StandardScaler()
df[numerical_columns]=ssr.fit_transform(df[numerical_columns])
```

```
[32]: df
```

```
[32]:
```

	Company	fueltype	aspiration	doornumber	carbody	drivewheel	\
0	1	1	0	1	0	2	
1	1	1	0	1	0	2	
2	1	1	0	1	2	2	
3	2	1	0	0	3	1	
4	2	1	0	0	3	0	
..	
200	26	1	0	0	3	2	
201	26	1	1	0	3	2	
202	26	1	0	0	3	2	
203	26	0	1	0	3	2	
204	26	1	1	0	3	2	

	enginelocation	carlength	carwidth	carheight	...	enginesize	\
0	0	-0.426521	-0.844782	-2.020417	...	0.074449	
1	0	-0.426521	-0.844782	-2.020417	...	0.074449	
2	0	-0.231513	-0.190566	-0.543527	...	0.604046	
3	0	0.207256	0.136542	0.235942	...	-0.431076	
4	0	0.207256	0.230001	0.235942	...	0.218885	
..	
200	0	1.198549	1.398245	0.728239	...	0.339248	
201	0	1.198549	1.351515	0.728239	...	0.339248	
202	0	1.198549	1.398245	0.728239	...	1.109571	
203	0	1.198549	1.398245	0.728239	...	0.435538	
204	0	1.198549	1.398245	0.728239	...	0.339248	

	fuelsystem	boreratio	stroke	compressionratio	horsepower	peakrpm	\
0	5	0.519071	-1.839377	-0.288349	0.174483	-0.262960	
1	5	0.519071	-1.839377	-0.288349	0.174483	-0.262960	
2	5	-2.404880	0.685946	-0.288349	1.264536	-0.262960	
3	5	-0.517266	0.462183	-0.035973	-0.053668	0.787855	
4	5	-0.517266	0.462183	-0.540725	0.275883	0.787855	
..	
200	5	1.666445	-0.336970	-0.162161	0.250533	0.577692	

201	5	1.666445	-0.336970	-0.364062	1.416637	0.367529
202	5	0.926204	-1.232021	-0.338824	0.757535	0.787855
203	3	-1.183483	0.462183	3.244916	0.047732	-0.683286
204	5	1.666445	-0.336970	-0.162161	0.250533	0.577692

	citympg	highwaympg	price
0	-0.646553	-0.546059	13495.0
1	-0.646553	-0.546059	16500.0
2	-0.953012	-0.691627	16500.0
3	-0.186865	-0.109354	13950.0
4	-1.106241	-1.273900	17450.0
..
200	-0.340094	-0.400490	16845.0
201	-0.953012	-0.837195	19045.0
202	-1.106241	-1.128332	21485.0
203	0.119594	-0.546059	22470.0
204	-0.953012	-0.837195	22625.0

[205 rows x 23 columns]

```
[33]: # x and y split
```

```
[34]: x=df.iloc[ : , :-1]
      y=df.iloc[ : , -1]
```

```
[35]: # Train and Test split
      from sklearn.model_selection import train_test_split
      x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.
      ↪3,random_state=11)
```

2 LinearRegression

```
[36]: # LinearRegression
      from sklearn.linear_model import LinearRegression

      # Create a Linear Regression model
      lr=LinearRegression()

      # Train the model on the training data
      lr.fit(x_train,y_train)

      # Make predictions on the test data
      y_pred_lr=lr.predict(x_test)

      # Evaluate the model's performance
      from sklearn.metrics import mean_absolute_error,mean_squared_error
```

```
print('MAE:',mean_absolute_error(y_test,y_pred_lr))
mse_lr=mean_squared_error(y_test,y_pred_lr)
print('MSE:',mse_lr)
```


MAE: 2142.9703979715828

MSE: 9351120.379315864

```
[37]: # Create a scatter plot
sns.scatterplot(x=y_test,y=y_pred_lr)

# Add labels and a title
plt.xlabel('Y_TEST')
plt.ylabel('Y_PREDICT_linear regression')
plt.title('EVALUATE')
```

```
[37]: Text(0.5, 1.0, 'EVALUATE')
```



3 Support Vector Machine (SVM)

```
[38]: #svm
from sklearn.svm import SVR

# Create a svm model
svr=SVR()

# Train the model on the training data
svr.fit(x_train,y_train)

# Make predictions on the test data
y_pred_svm=svr.predict(x_test)

# Evaluate the model
from sklearn.metrics import mean_absolute_error,mean_squared_error
print('MAE:',mean_absolute_error(y_test,y_pred_svm))
mse_svm=mean_squared_error(y_test,y_pred_svm)
print("mse",mse_svm)
```


MAE: 4427.266337086366

mse 48774553.675131686

```
[39]: # Create a scatter plot
sns.scatterplot(x=y_test,y=y_pred_svm,color = "red")

# Add labels and a title
plt.xlabel('Y_TEST')
plt.ylabel('Y_PREDICT_SVM')
plt.title('EVALUATE')
```

```
[39]: Text(0.5, 1.0, 'EVALUATE')
```



4 KNeighborsRegressor

```
[40]: # KNeighborsRegressor
from sklearn.neighbors import KNeighborsRegressor

## Create KNN model
k=3
knn=KNeighborsRegressor(n_neighbors=k)

# Train the model on the training data
knn.fit(x_train,y_train)

# Make predictions on the test data
y_pred_knn=knn.predict(x_test)

# Evaluate the model's performance
from sklearn.metrics import mean_absolute_error,mean_squared_error
print('MAE:',mean_absolute_error(y_test,y_pred_knn))
mse_knn=mean_squared_error(y_test,y_pred_knn)
```

```
print("MSE",mse_knn)
```


MAE: 1955.4408548387094

MSE 9697257.796603944

```
[41]: # Create a scatter plot
sns.scatterplot(x=y_test,y=y_pred_knn,color = "y")

# Add labels and a title
plt.xlabel('Y_TEST')
plt.ylabel('Y_PREDICT_KNN')
plt.title('EVALUATE')
```

```
[41]: Text(0.5, 1.0, 'EVALUATE')
```



```
[42]: #Choosing 'k' by elbow method
```

```
[43]: knn_mse=[]
for i in range(1,60):
    knn_i=KNeighborsRegressor(n_neighbors=i)
```

```
# train data
knn_i.fit(x_train,y_train)
# predict
y_pred_knn_i=knn.predict(x_test)
knn_mse.append(mean_squared_error(y_test,y_pred_knn_i))
```

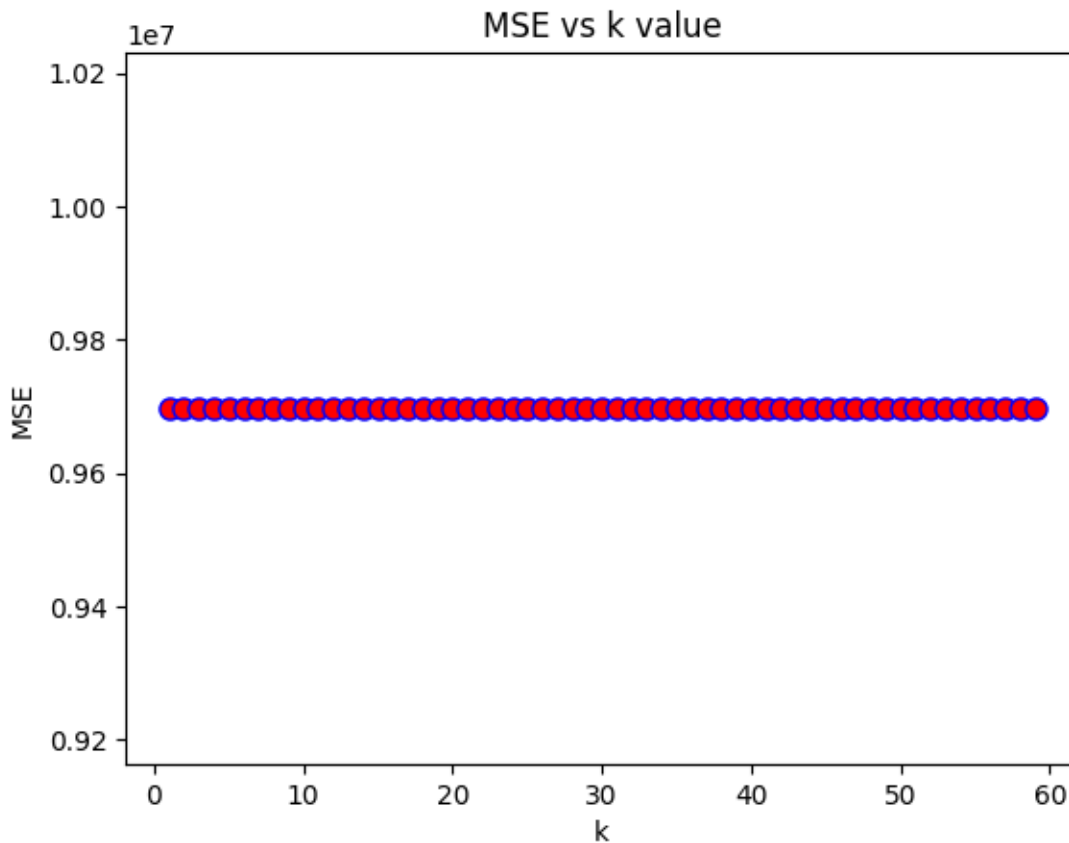
knn_mse

[illegible]

[illegible]

```
[45]: # PLOT FOR ERROR VALUE
plt.
    plot(range(1,60),knn_mse,color='blue',linestyle='dashed',marker='o',markerfacecolor='red',
    markersize=8)
plt.title('MSE vs k value')
plt.xlabel('k')
plt.ylabel("MSE")
```

```
[45]: Text(0, 0.5, 'MSE')
```



```
[46]: #DecisionTreeRegressor
```

```
[47]: # DecisionTreeRegressor
```

```
from sklearn.tree import DecisionTreeRegressor

# Create a # Create a Linear Regression model model
dt=DecisionTreeRegressor()

# Train the model on the training data
dt.fit(x_train,y_train)

# Make predictions on the test data
y_pred_dt=dt.predict(x_test)


# Evaluate the model's performance
print('MAE:',mean_absolute_error(y_test,y_pred_dt))
mse_dt=mean_squared_error(y_test,y_pred_dt)
print('MSE:',mse_dt)
```

MAE: 1700.9166612903225

MSE: 6275232.620449822


```
[48]: # Create a scatter plot
sns.scatterplot(x=y_test,y=y_pred_dt,color = "green")
# Add labels and a title
plt.xlabel('Y_TEST')
plt.ylabel('Y_PREDICT_DecisionTreeRegressor')
plt.title('EVALUATE')
```

```
[48]: Text(0.5, 1.0, 'EVALUATE')
```



```
[52]: class_name = ('linear regression','svm','knn','decision tree')
class_score=(mse_lr,mse_svm,mse_knn,mse_dt)
colors=('r','g','orange','b','pink')
plt.bar(class_name,class_score,color=colors)
plt.title ("mean_squared_error ofthe classes ")
plt.ylabel('mse')
```

```
[52]: Text(0, 0.5, 'mse')
```



5 r2_score

```
[53]: from sklearn.metrics import r2_score
```

```
[54]: r2_lr=r2_score(y_test,y_pred_lr)
      print('r2_score_lr',r2_lr)
      r2_svm=r2_score(y_test,y_pred_svm)
      print('r2_score_svm',r2_svm)
      r2_knn=r2_score(y_test,y_pred_knn)
      print('r2_score_knn',r2_knn)
      r2_dt=r2_score(y_test,y_pred_dt)
      print('r2_score_knn',r2_dt)
```

```
r2_score_lr 0.7995224343756604
r2_score_svm -0.045671896902632625
r2_score_knn 0.7921016351586058
r2_score_knn 0.8654660288346897
```

```
[55]: #r2 in function
```



```
[56]: def r2score(Y_test,Y_pred):
      print('R2 SCORE',r2_score(Y_test,Y_pred))
```

```
[57]: #LR
      r2score(y_test,y_pred_lr)
      #KNN
      r2score(y_test,y_pred_knn)
      # SVM
      r2score(y_test,y_pred_svm)
      # DT
      r2score(y_test,y_pred_dt)
```

```
R2 SCORE 0.7995224343756604
R2 SCORE 0.7921016351586058
R2 SCORE -0.045671896902632625
R2 SCORE 0.8654660288346897
```

```
[58]: #data frame
```

```
[59]: data={"model":['LinearRegression','SVM','KNN','DecisionTreeRegressor'],
          "mse": [mse_lr,mse_svm,mse_knn,mse_dt],
          'r2score': [r2_lr,r2_svm,r2_knn,r2_dt]}
      df2=pd.DataFrame(data)
      df2.T
```

```
[59]:
```

	0	1	2 \
model	LinearRegression	SVM	KNN
mse	9351120.379316	48774553.675132	9697257.796604
r2score	0.799522	-0.045672	0.792102

	3
model	DecisionTreeRegressor
mse	6275232.62045
r2score	0.865466